

```
In [24]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import math
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

```
In [25]: df=pd.read_csv('BostonHousing.csv')
```

```
In [26]: df.info()
# count This shows the number of non-null values in each numerical column
#mean
#std ,
#25%= Q1 value below which 25% of the data falls.
# 50% median
#75%= Q3 value below which 75% of the data falls.
#max and min val in that column
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   crim        506 non-null    float64
1   zn           506 non-null    float64
2   indus        506 non-null    float64
3   chas         506 non-null    int64
4   nox          506 non-null    float64
5   rm           501 non-null    float64
6   age          506 non-null    float64
7   dis          506 non-null    float64
8   rad          506 non-null    int64
9   tax          506 non-null    int64
10  ptratio      506 non-null    float64
11  b            506 non-null    float64
12  lstat        506 non-null    float64
13  medv         506 non-null    float64
dtypes: float64(11), int64(3)
memory usage: 55.5 KB
```

```
In [27]: df.describe()
```

Out [27]:

	crim	zn	indus	chas	nox	rm	
count	506.000000	506.000000	506.000000	506.000000	506.000000	501.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284341	6.284341
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.705587	0.705587
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.884000	5.884000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208000	6.208000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.625000	6.625000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000

In [28]: `df.head()`

Out [28]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90

In [29]: `df.shape`

Out [29]: (506, 14)

In [30]: `df.tail()`

Out [30]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	21.0	396.90

In [31]: `df.corr()`

Out [31]:

	crim	zn	indus	chas	nox	rm	age
crim	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219433	0.352734
zn	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311173	-0.569537
indus	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.394193	0.644779
chas	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091468	0.086518
nox	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302751	0.731470
rm	-0.219433	0.311173	-0.394193	0.091468	-0.302751	1.000000	-0.240286
age	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240286	1.000000
dis	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.203507	-0.747881
rad	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.210718	0.456022
tax	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292794	0.506456
ptratio	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.357612	0.261515
b	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128107	-0.273534
lstat	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.615721	0.602339
medv	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.696169	-0.376955

In [32]: `df.isnull()`

Out [32]:

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b	l
0	False	False	False	False	False	False	False	False	False	False	False	False	F
1	False	False	False	False	False	False	False	False	False	False	False	False	F
2	False	False	False	False	False	False	False	False	False	False	False	False	F
3	False	False	False	False	False	False	False	False	False	False	False	False	F
4	False	False	False	False	False	False	False	False	False	False	False	False	F
...
501	False	False	False	False	False	False	False	False	False	False	False	False	F
502	False	False	False	False	False	False	False	False	False	False	False	False	F
503	False	False	False	False	False	False	False	False	False	False	False	False	F
504	False	False	False	False	False	False	False	False	False	False	False	False	F
505	False	False	False	False	False	False	False	False	False	False	False	False	F

506 rows × 14 columns

In [33]: `df.isnull().sum()`

```
Out[33]: crim      0
          zn        0
          indus     0
          chas      0
          nox       0
          rm        5
          age       0
          dis       0
          rad       0
          tax       0
          ptratio   0
          b         0
          lstat     0
          medv      0
          dtype: int64
```

```
In [34]: df['rm']=df['rm'].fillna(df['rm'].mean())
```

```
In [35]: df.isnull().sum()
```

```
Out[35]: crim      0
          zn        0
          indus     0
          chas      0
          nox       0
          rm        0
          age       0
          dis       0
          rad       0
          tax       0
          ptratio   0
          b         0
          lstat     0
          medv      0
          dtype: int64
```

```
In [36]: df.head()
```

```
Out[36]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	b
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90

```
In [37]: X=df.drop('medv',axis=1)
          y=df['medv']
          df.dropna(inplace=True)
          print(X)
          print(y)
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	ptratio	b	lstat
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33
..
501	21.0	391.99	9.67
502	21.0	396.90	9.08
503	21.0	396.90	5.64
504	21.0	393.45	6.48
505	21.0	396.90	7.88

[506 rows x 13 columns]

0	24.0
1	21.6
2	34.7
3	33.4
4	36.2

..	...
501	22.4
502	20.6
503	23.9
504	22.0
505	11.9

Name: medv, Length: 506, dtype: float64

In [38]: `X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_stat`

In [39]: `model=LinearRegression()
model.fit(X_train,y_train)`

Out[39]: `LinearRegression
LinearRegression()`


In [40]: `print("Shape of X_train: ",X_train.shape)
print("Shape of X_test: ", X_test.shape)
print("Shape of y_train: ",y_train.shape)
print("Shape of y_test",y_test.shape)`

```
Shape of X_train: (404, 13)
Shape of X_test: (102, 13)
Shape of y_train: (404,)
Shape of y_test (102,)
```

```
In [41]: pred = model.predict(X_test)
```

```
In [42]: plt.figure(figsize=(6,6));
plt.scatter(y_test,pred);
plt.xlabel('Y test')
plt.ylabel('Predicted Y')
plt.title('Test vs Predict')
```

```
Out[42]: Text(0.5, 1.0, 'Test vs Predict')
```




```
In [43]: from sklearn import metrics

print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
```

Mean Absolute Error: 3.2064039639004025
Mean Squared Error: 24.404825188146653
Root Mean Squared Error: 4.940124005341025

```
In [44]: import matplotlib.pyplot as plt

plt.figure(figsize=(8, 6)) # Adjust figure size for better visibility
plt.scatter(y_test, pred, alpha=0.5, color='blue') # Add transparency and s
plt.xlabel('Actual Test Values', fontsize=12) # Customize x-axis label
plt.ylabel('Predicted Test Values', fontsize=12) # Customize y-axis label
plt.title('Actual vs. Predicted Test Values', fontsize=14) # Customize plot
plt.grid(True) # Add grid lines
plt.show()
```



```
In [46]: r2 = r2_score(y_test, pred)
print('R-squared Score:', r2)
```

R-squared Score: 0.6672089705941856

In []: